

## hw2

July 8, 2021

```
[187]: # !pip install imblearn
```

```
Collecting imblearn
  Downloading imblearn-0.0-py2.py3-none-any.whl (1.9 kB)
Collecting imbalanced-learn
  Downloading imbalanced_learn-0.8.0-py3-none-any.whl (206 kB)
Requirement already satisfied: joblib>=0.11 in
c:\users\jack\miniconda3\envs\datascience\lib\site-packages (from imbalanced-learn->imblearn) (1.0.1)
Requirement already satisfied: scikit-learn>=0.24 in
c:\users\jack\miniconda3\envs\datascience\lib\site-packages (from imbalanced-learn->imblearn) (0.24.2)
Requirement already satisfied: numpy>=1.13.3 in
c:\users\jack\miniconda3\envs\datascience\lib\site-packages (from imbalanced-learn->imblearn) (1.19.2)
Requirement already satisfied: scipy>=0.19.1 in
c:\users\jack\miniconda3\envs\datascience\lib\site-packages (from imbalanced-learn->imblearn) (1.6.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
c:\users\jack\miniconda3\envs\datascience\lib\site-packages (from scikit-learn>=0.24->imbalanced-learn->imblearn) (2.1.0)
Installing collected packages: imbalanced-learn, imblearn
Successfully installed imbalanced-learn-0.8.0 imblearn-0.0
```

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import lightgbm as lgb
from collections import Counter
from matplotlib import pyplot as plt
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.metrics import roc_auc_score, accuracy_score

%matplotlib inline
```

```
[2]: train_data = pd.read_csv(r'train_final.csv')
test_data = pd.read_csv(r'test_final.csv')
```

```
50000  146      loan_status
```

```
[3]: print("train data size: ", train_data.shape)
print("test data size: ", test_data.shape)
```

```
train data size: (50000, 146)
```

```
test data size: (50000, 146)
```

```
[4]: train_data['loan_status'].head()
```

```
[4]: 0    1
1    1
2    1
3    1
4    1
Name: loan_status, dtype: int64
```

```
[5]: train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50000 entries, 0 to 49999
Columns: 146 entries, continuous_annual_inc to discrete_term_2_one_hot
dtypes: float64(20), int64(126)
memory usage: 55.7 MB
```

```
[77]: # [idx for idx in train_data.columns if train_data[idx].dtype == np.float64]
```

## 0.1

```
[7]: train_data.isnull().sum().sort_values(ascending=False)[:10]
```

```
[7]: continuous_annual_inc_joint      49780
continuous_dti_joint                49780
continuous_mths_since_last_record    40505
continuous_mths_since_last_major_derog 34948
continuous_mths_since_last_delinq     23917
continuous_dti                       1
discrete_addr_state_29_one_hot        0
discrete_addr_state_34_one_hot        0
discrete_addr_state_33_one_hot        0
discrete_addr_state_32_one_hot        0
dtype: int64
```

```
[8]: sum(train_data['continuous_annual_inc_joint'].isnull() &
↳ train_data['loan_status'] == 1) /
↳ sum(train_data['continuous_annual_inc_joint'].isnull()),\
```

```
sum(train_data['continuous_annual_inc_joint'].notnull() &
↳train_data['loan_status'] == 1) /
↳sum(train_data['continuous_annual_inc_joint'].notnull())
```

[8]: (0.7960024106066693, 0.740909090909091)

```
[9]: sum(train_data['continuous_dti_joint'].isnull() & train_data['loan_status'] ==
↳1)/sum(train_data['continuous_dti_joint'].isnull()),\
sum(train_data['continuous_dti_joint'].notnull() & train_data['loan_status'] ==
↳1)/sum(train_data['continuous_dti_joint'].notnull())
```

[9]: (0.7960024106066693, 0.740909090909091)

```
[10]: record = sum(train_data['continuous_mths_since_last_record'].isnull())
TP = sum(train_data['continuous_mths_since_last_record'].isnull() &
↳train_data['loan_status'] == 1)
precision = TP / sum(train_data['continuous_mths_since_last_record'].isnull())
recall = TP / sum(train_data['loan_status'] == 1)
2 * precision * recall / (precision + recall)
```

[10]: 0.8095848953208872

```
[11]: sum(train_data['continuous_mths_since_last_record'].isnull() &
↳train_data['loan_status'] == 1) /
↳sum(train_data['continuous_mths_since_last_record'].isnull()),\
sum(train_data['continuous_mths_since_last_record'].notnull() &
↳train_data['loan_status'] == 1) /
↳sum(train_data['continuous_mths_since_last_record'].notnull())
```


[11]: (0.8024194543883472, 0.7673512374934176)

```
[12]: major_derog = sum(train_data['continuous_mths_since_last_major_derog'].isnull())
major_derog_loan = sum(train_data['continuous_mths_since_last_major_derog'].
↳isnull() & train_data['loan_status'] == 1)
major_derog, major_derog_loan, major_derog_loan / major_derog
```

[12]: (34948, 28055, 0.802764106672771)

```
[13]: sns.histplot(train_data['continuous_mths_since_last_major_derog'])
```

[13]: <AxesSubplot:xlabel='continuous\_mths\_since\_last\_major\_derog', ylabel='Count'>




```
[14]: delinq = sum(train_data['continuous_mths_since_last_delinq'].isnull())
delinq_loan = sum(train_data['continuous_mths_since_last_delinq'].isnull() &
    ↳ train_data['loan_status'] == 1)
delinq, delinq_loan, delinq_loan / delinq
```

```
[14]: (23917, 19200, 0.8027762679265794)
```

```
[15]: sns.histplot(train_data['continuous_mths_since_last_delinq'])
```

```
[15]: <AxesSubplot:xlabel='continuous_mths_since_last_delinq', ylabel='Count'>
```



0.2

- continuous\_dti
- continuous\_annual\_inc\_joint continuous\_dti\_joint null
- null binary

```
[16]: train_data.fillna(value={'continuous_dti': train_data['continuous_dti'].
    ↳mean()}, inplace=True)
train_data['continuous_mths_since_last_record_isnull'] =
    ↳train_data['continuous_mths_since_last_record'].isnull()
train_data['continuous_mths_since_last_major_derog_isnull'] =
    ↳train_data['continuous_mths_since_last_major_derog'].isnull()
train_data['continuous_mths_since_last_delinq_isnull'] =
    ↳train_data['continuous_mths_since_last_delinq'].isnull()
train_data.drop(columns=['continuous_annual_inc_joint',
    'continuous_dti_joint',
    'continuous_mths_since_last_record',
    'continuous_mths_since_last_major_derog',
    'continuous_mths_since_last_delinq'], inplace=True)
```

```
[17]: test_data.fillna(value={'continuous_dti': test_data['continuous_dti'].mean()},
    ↳inplace=True)
test_data['continuous_mths_since_last_record_isnull'] =
    ↳test_data['continuous_mths_since_last_record'].isnull()
```

```
test_data['continuous_mths_since_last_major_derog_isnull'] =_
↳test_data['continuous_mths_since_last_major_derog'].isnull()
test_data['continuous_mths_since_last_delinq_isnull'] =_
↳test_data['continuous_mths_since_last_delinq'].isnull()
test_data.drop(columns=['continuous_annual_inc_joint',
                        'continuous_dti_joint',
                        'continuous_mths_since_last_record',
                        'continuous_mths_since_last_major_derog',
                        'continuous_mths_since_last_delinq'], inplace=True)
```

```
[18]: train_data.isnull().sum().sort_values(ascending=False)[:5]
```

```
[18]: continuous_mths_since_last_delinq_isnull      0
continuous_mths_since_last_major_derog_isnull    0
discrete_addr_state_36_one_hot                   0
discrete_addr_state_35_one_hot                   0
discrete_addr_state_34_one_hot                   0
dtype: int64
```

```
[19]: train_data.columns
```


```
[19]: Index(['continuous_annual_inc', 'continuous_delinq_2yrs', 'continuous_dti',
          'continuous_fico_range_high', 'continuous_fico_range_low',
          'continuous_funded_amnt', 'continuous_funded_amnt_inv',
          'continuous_inq_last_6mths', 'continuous_installment',
          'continuous_int_rate',
          ...,
          'discrete_sub_grade_31_one_hot', 'discrete_sub_grade_32_one_hot',
          'discrete_sub_grade_33_one_hot', 'discrete_sub_grade_34_one_hot',
          'discrete_sub_grade_35_one_hot', 'discrete_term_1_one_hot',
          'discrete_term_2_one_hot', 'continuous_mths_since_last_record_isnull',
          'continuous_mths_since_last_major_derog_isnull',
          'continuous_mths_since_last_delinq_isnull'],
          dtype='object', length=144)
```

### 0.3

```
loan_status 0    10000    1    40000          ( SMOTE          )
```


```
[20]: sns.countplot(x='loan_status', data=train_data)
```

```
[20]: <AxesSubplot:xlabel='loan_status', ylabel='count'>
```



```
[21]: sns.countplot(x='loan_status', data=test_data)
```

```
[21]: <AxesSubplot:xlabel='loan_status', ylabel='count'>
```



## 0.4 one hot

lightgbm

LightGBM can use categorical features as input directly. It doesn't need to convert to one-hot

```
[22]: # one_hots = [idx for idx in train_data.columns if 'one_hot' in idx]
# one_hot_dict = {}
# for string in one_hots:
#     string = string[:-8]
#     n = int(string.split('_')[-1])
#     s = '_'.join(string.split('_')[:-1])
#     if s not in one_hot_dict.keys():
#         one_hot_dict[s] = n
#     else:
#         one_hot_dict[s] = max(one_hot_dict[s], n)
# one_hot_dict
```

```
[22]: {'discrete_addr_state': 49,
'discrete_application_type': 2,
'discrete_emp_length': 12,
'discrete_grade': 7,
'discrete_home_ownership': 4,
'discrete_policy_code': 1,
'discrete_purpose': 12,
'discrete_pymnt_plan': 1,
'discrete_sub_grade': 35,
'discrete_term': 2}
```

```
[105]: # def helperOneHot(train_data, s, n):
#     for i in range(train_data.shape[0]):
#         for k in range(1, n + 1):
#             if train_data.loc[i, '{}{}_one_hot'.format(s, k)] == 1:
#                 train_data.loc[i, s] = int(k)
#                 break
#     return train_data.astype({s: 'int64'})

# train_data_no_oh = train_data
# test_data_no_oh = test_data
# for s, n in one_hot_dict.items():
#     train_data_no_oh = helperOneHot(train_data_no_oh, s, n)
#     test_data_no_oh = helperOneHot(test_data_no_oh, s, n)

# for s in one_hot_dict.keys():
```



```
#     cols = ['{}_{}_one_hot'.format(s, i) for i in range(1, one_hot_dict[s] +
→1)]
#     train_data_no_oh.drop(columns=cols, inplace=True)
#     test_data_no_oh.drop(columns=cols, inplace=True)
```

## 0.5

```
[106]: # X_train_no_oh, y_train_no_oh = train_data_no_oh.drop(columns='loan_status'),
→train_data_no_oh['loan_status']
# print(X_train_no_oh.shape, y_train_no_oh.shape)

# X_test_no_oh, y_test_no_oh = test_data_no_oh.drop(columns='loan_status'),
→test_data_no_oh['loan_status']
# print(X_test_no_oh.shape, y_test_no_oh.shape)
```

(50000, 28) (50000,)

(50000, 28) (50000,)

```
[107]: # counter = Counter(y_train)
# print(counter)
# # # transform the dataset
# oversample = SMOTE()
# X_train_smote, y_train_smote = oversample.fit_resample(X_train_no_oh,
→y_train_no_oh)
# counter = Counter(y_train_smote)
# print(counter)
```

Counter({0: 20387, 1: 20358})

Counter({1: 39788, 0: 39788})

lgb.cv    GridSearchCV

```
[78]: # X_train, X_val, y_train, y_val = train_test_split(X_train_no_oh,
→y_train_no_oh, test_size=0.1, stratify=y_train_no_oh)
# X_train, X_val, y_train, y_val = train_test_split(X_train_no_oh,
→y_train_no_oh, test_size=0.1)
# X_train, X_val, y_train, y_val = train_test_split(train_data.
→drop(columns='loan_status'), train_data['loan_status'], test_size=0.1)
# X_train, X_val, y_train, y_val = train_test_split(X_train_smote,
→y_train_smote, test_size=0.2, stratify=y_train_smote)

# print(X_train.head())
print(X_train.shape, X_val.shape)
```

(45000, 143) (5000, 143)

lgb.Dataset

```
[53]: # lgb.Dataset
# train_lgb = lgb.Dataset(X_train, label=y_train,
# →categorical_feature=list(one_hot_dict.keys()), free_raw_data=False)
# validation_lgb = lgb.Dataset(X_val, label=y_val,
# →categorical_feature=list(one_hot_dict.keys()), free_raw_data=False)
data_lgb = lgb.Dataset(train_data.drop(columns='loan_status'),
# →label=train_data['loan_status'], free_raw_data=False)
train_lgb = lgb.Dataset(X_train, label=y_train, free_raw_data=False)
validation_lgb = lgb.Dataset(X_val, label=y_val, free_raw_data=False)

#
params = {
    'boosting_type': 'gbdt',
    'objective': 'binary',
    'learning_rate': 0.1,
    'metric': 'auc',
    'num_iterations': 1000,
    'num_leaves': 20,
    'max_depth': 4,
    'subsample': 0.8,
    'colsample_bytree': 0.8
}
```

```
lgb.cv      n_estimators    83
```

```
[54]: #
# gbm = lgb.train(params, train_lgb, valid_sets=[validation_lgb])

cv_results = lgb.cv(params, data_lgb, num_boost_round=1000, nfold=5,
# →stratified=False, shuffle=True, metrics='auc', early_stopping_rounds=50,
# →seed=0)
print('best n_estimators:', len(cv_results['auc-mean']))
print('best cv score:', pd.Series(cv_results['auc-mean']).max())
```

C:\Users\Jack\miniconda3\envs\datascience\lib\site-packages\lightgbm\engine.py:527: UserWarning: Found `num\_iterations` in params. Will use it instead of argument  
\_log\_warning("Found `{}` in params. Will use it instead of argument".format(alias))

[LightGBM] [Info] Number of positive: 31856, number of negative: 8144  
[LightGBM] [Warning] Auto-choosing row-wise multi-threading, the overhead of testing was 0.000933 seconds.  
You can set `force\_row\_wise=true` to remove the overhead.  
And if memory is not enough, you can set `force\_col\_wise=true`.  
[LightGBM] [Info] Total Bins 2152  
[LightGBM] [Info] Number of data points in the train set: 40000, number of used

```

features: 139
[LightGBM] [Info] Number of positive: 31852, number of negative: 8148
[LightGBM] [Warning] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.000876 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2152
[LightGBM] [Info] Number of data points in the train set: 40000, number of used
features: 139
[LightGBM] [Info] Number of positive: 31795, number of negative: 8205
[LightGBM] [Warning] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.000933 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2152
[LightGBM] [Info] Number of data points in the train set: 40000, number of used
features: 139
[LightGBM] [Info] Number of positive: 31798, number of negative: 8202
[LightGBM] [Warning] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.000831 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2152
[LightGBM] [Info] Number of data points in the train set: 40000, number of used
features: 139
[LightGBM] [Info] Number of positive: 31851, number of negative: 8149
[LightGBM] [Warning] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.000895 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 2152
[LightGBM] [Info] Number of data points in the train set: 40000, number of used
features: 139
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.796400 -> initscore=1.363944
[LightGBM] [Info] Start training from score 1.363944
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.796300 -> initscore=1.363328
[LightGBM] [Info] Start training from score 1.363328
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.794875 -> initscore=1.354565
[LightGBM] [Info] Start training from score 1.354565
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.794950 -> initscore=1.355025
[LightGBM] [Info] Start training from score 1.355025
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.796275 -> initscore=1.363174
[LightGBM] [Info] Start training from score 1.363174
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf
[LightGBM] [Warning] No further splits with positive gain, best gain: -inf

```

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]



[illegible]

[illegible]





```
[57]: clf1 = gsearch1.best_estimator_
      clf1
```

```
[57]: LGBMClassifier(bagging_fraction=0.8, feature_fraction=0.8, max_depth=4,
                    metrics='auc', n_estimators=83, num_leaves=10,
                    objective='binary')
```

optimal value for max\_depth is 4 and optimal value for num\_leaves is 10.

```
accuracy 0.91764 roc_auc 0.875699
```

```
[58]: def results(clf, data):
      y_pred = clf.predict(data.drop(columns='loan_status'))
      y_pred = np.array(list(map(lambda x: 1 if x >= 0.5 else 0, y_pred)))
      y = data['loan_status']
      return {'accuracy': accuracy_score(y, y_pred),
              'roc_auc': roc_auc_score(y, y_pred)}

      results(clf1, test_data)
```

```
[58]: {'accuracy': 0.91764, 'roc_auc': 0.8756988382937406}

      min_data_in_leaf 140
```

```
[59]: paramTest2 = {'min_data_in_leaf': range(5, 155, 5)}

      gsearch2 = GridSearchCV(estimator = lgb.LGBMClassifier(boosting_type='gbdt',
                                                            objective='binary',
                                                            metrics='auc',
                                                            learning_rate=0.1,
                                                            n_estimators=82,
                                                            max_depth=4,
                                                            num_leaves=10,
                                                            bagging_fraction = 0.8,
                                                            feature_fraction = 0.8),
                             param_grid=paramTest2, scoring='roc_auc', cv=5,
                             n_jobs=-1)
      gsearch2.fit(X_train, y_train)
      gsearch2.cv_results_['mean_test_score'], gsearch2.
      cv_results_['mean_test_score'].mean()
```

[LightGBM] [Warning] feature\_fraction is set=0.8, colsample\_bytree=1.0 will be ignored. Current value: feature\_fraction=0.8

[LightGBM] [Warning] min\_data\_in\_leaf is set=140, min\_child\_samples=20 will be ignored. Current value: min\_data\_in\_leaf=140

[LightGBM] [Warning] bagging\_fraction is set=0.8, subsample=1.0 will be ignored. Current value: bagging\_fraction=0.8

```
[59]: (array([0.96059155, 0.96057467, 0.96053096, 0.96057324, 0.96055448,
             0.96055347, 0.96061442, 0.96056771, 0.9606309 , 0.96072561,
             0.96068197, 0.96066348, 0.96071941, 0.9607305 , 0.96061674,
             0.96058627, 0.96060219, 0.96053595, 0.96059588, 0.96066609,
             0.9606667 , 0.96071893, 0.96077546, 0.96059836, 0.96064973,
             0.96074422, 0.96067689, 0.96078756, 0.96063708, 0.96052441]),
      0.9606364933221986)
```

```
[60]: gsearch2.best_params_
```

```
[60]: {'min_data_in_leaf': 140}
```

```
[61]: clf2 = gsearch2.best_estimator_
      clf2
```

```
[61]: LGBMClassifier(bagging_fraction=0.8, feature_fraction=0.8, max_depth=4,
                    metrics='auc', min_data_in_leaf=140, n_estimators=82,
                    num_leaves=10, objective='binary')
```

```
[62]: results(clf2, test_data)
```

```
[62]: {'accuracy': 0.91762, 'roc_auc': 0.8753765976682328}
```

```
lambda_l1 lambda_l2
```

```
[63]: paramTest3 = {'lambda_l1': [1e-5, 1e-3, 1e-1, 0.0, 0.1, 0.3, 0.5, 0.7, 0.9, 1.
    ↪0],
                   'lambda_l2': [1e-5, 1e-3, 1e-1, 0.0, 0.1, 0.3, 0.5, 0.7, 0.9, 1.0]}

gsearch3 = GridSearchCV(estimator = lgb.LGBMClassifier(boosting_type='gbdt',
                                                         objective='binary',
                                                         metrics='auc',
                                                         learning_rate=0.1,
                                                         n_estimators=82,
                                                         min_data_in_leaf=140,
                                                         max_depth=4,
                                                         num_leaves=10,
                                                         bagging_fraction = 0.8,
                                                         feature_fraction = 0.8),
                       param_grid=paramTest3, scoring='roc_auc', cv=5,
    ↪n_jobs=-1)
gsearch3.fit(X_train,y_train)
gsearch3.cv_results_['mean_test_score'], gsearch3.
    ↪cv_results_['mean_test_score'].mean()
```

[LightGBM] [Warning] feature\_fraction is set=0.8, colsample\_bytree=1.0 will be

```

ignored. Current value: feature_fraction=0.8
[LightGBM] [Warning] min_data_in_leaf is set=140, min_child_samples=20 will be
ignored. Current value: min_data_in_leaf=140
[LightGBM] [Warning] lambda_l1 is set=0.001, reg_alpha=0.0 will be ignored.
Current value: lambda_l1=0.001
[LightGBM] [Warning] bagging_fraction is set=0.8, subsample=1.0 will be ignored.
Current value: bagging_fraction=0.8
[LightGBM] [Warning] lambda_l2 is set=0.001, reg_lambda=0.0 will be ignored.
Current value: lambda_l2=0.001

```

```

[63]: (array([0.96078755, 0.96078753, 0.96062171, 0.96078756, 0.96062171,
              0.96062813, 0.96063375, 0.9606609 , 0.96061072, 0.96056447,
              0.96077298, 0.96080488, 0.96060907, 0.96077296, 0.96060907,
              0.96062493, 0.96063252, 0.96067345, 0.96059741, 0.9605999 ,
              0.96063565, 0.96061885, 0.96070876, 0.96063565, 0.96070876,
              0.96069257, 0.96062434, 0.96061176, 0.96060351, 0.96067026,
              0.96078756, 0.96078753, 0.96062171, 0.96078756, 0.96062171,
              0.96062813, 0.96063374, 0.96066087, 0.96061072, 0.96056447,
              0.96063565, 0.96061885, 0.96070876, 0.96063565, 0.96070876,
              0.96069257, 0.96062434, 0.96061176, 0.96060351, 0.96067026,
              0.96065666, 0.96065662, 0.96069431, 0.96065666, 0.96069431,
              0.96055566, 0.96072729, 0.96076221, 0.96072778, 0.9606771 ,
              0.96061636, 0.96061637, 0.96061122, 0.96061634, 0.96061122,
              0.96058555, 0.96070384, 0.96062638, 0.96061826, 0.96062976,
              0.96076249, 0.96076236, 0.96075459, 0.96076249, 0.96075459,
              0.96063745, 0.96065886, 0.96065782, 0.96059195, 0.96060302,
              0.96062758, 0.96062766, 0.96067657, 0.96062758, 0.96067657,
              0.96060148, 0.9606397 , 0.96058609, 0.96073001, 0.96075929,
              0.96069248, 0.96069239, 0.96061304, 0.96069248, 0.96061304,
              0.96064063, 0.96062726, 0.96068497, 0.96068147, 0.96067488]),
       0.9606628013226058)

```

```

[64]: gsearch3.best_params_

```

```

[64]: {'lambda_l1': 0.001, 'lambda_l2': 0.001}

```

```

[65]: clf3 = gsearch3.best_estimator_
      clf3

```

```

[65]: LGBMClassifier(bagging_fraction=0.8, feature_fraction=0.8, lambda_l1=0.001,
                    lambda_l2=0.001, max_depth=4, metrics='auc',
                    min_data_in_leaf=140, n_estimators=82, num_leaves=10,
                    objective='binary')

```

```

[66]: results(clf3, test_data)

```

```

[66]: {'accuracy': 0.91762, 'roc_auc': 0.8753765976682328}

```

```

learning_rate    0.01 n_estimators    800

```

```
[76]: clf = lgb.LGBMClassifier(
        learning_rate=0.01,
        n_estimators=800,
        bagging_fraction=0.8,
        feature_fraction=0.8,
        #     lambda_l1=0.1,
        #     lambda_l2=0.5,
        max_depth=4,
        boosting_type='gbdt',
        metrics='auc',
        min_data_in_leaf=140,
        num_leaves=10,
        objective='binary'
    )
    clf.fit(X_train, y_train)
    results(clf, test_data)
```

```
[LightGBM] [Warning] feature_fraction is set=0.8, colsample_bytree=1.0 will be
ignored. Current value: feature_fraction=0.8
[LightGBM] [Warning] min_data_in_leaf is set=140, min_child_samples=20 will be
ignored. Current value: min_data_in_leaf=140
[LightGBM] [Warning] bagging_fraction is set=0.8, subsample=1.0 will be ignored.
Current value: bagging_fraction=0.8
```

```
[76]: {'accuracy': 0.91788, 'roc_auc': 0.8753445529179209}
```

```
accuracy 0.91788 roc_auc 0.875345
```

```
[ ]:
```